

Indexing for Near-Sorted Data

Aneesh Raman

Subhadeep Sarkar


Matthaios Olma

Manos Athanassoulis


BOSTON
UNIVERSITY




Indexes in Databases




organize
data



efficient
queries




unstructured
data




structured
data



The process of inducing “*sortedness*” to an otherwise unsorted data collection




What if data already
has some structure?





What if data already
has some structure?




Near-sorted data



What if data already
has some structure?




\approx



Near-sorted data

treated same as
unstructured data!


Intermediate-Sortedness in Practice




Time Series




Stock market



Join/query



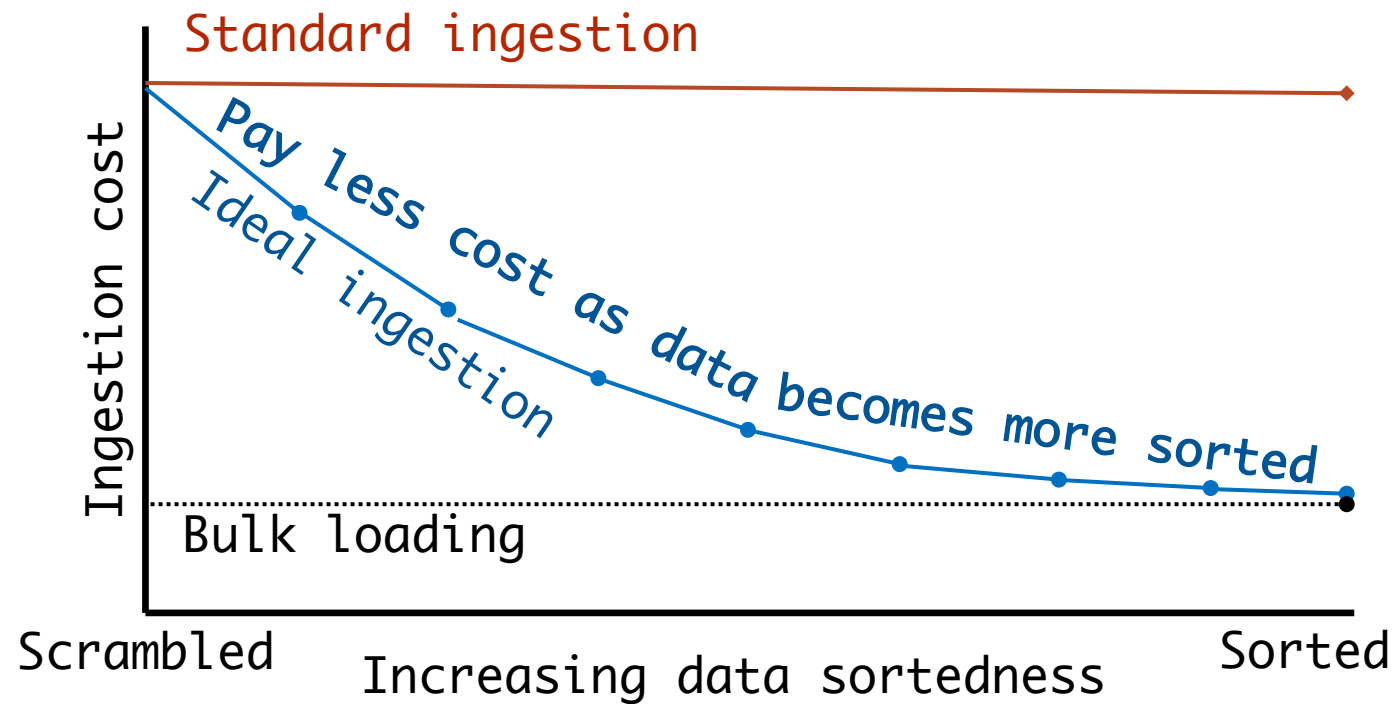
efficient reads



fast writes


classical indexes carry
redundant effort!

Ideally...




The Sortedness-Aware (SWARE) Paradigm

Sortedness-Aware (SWARE) Paradigm

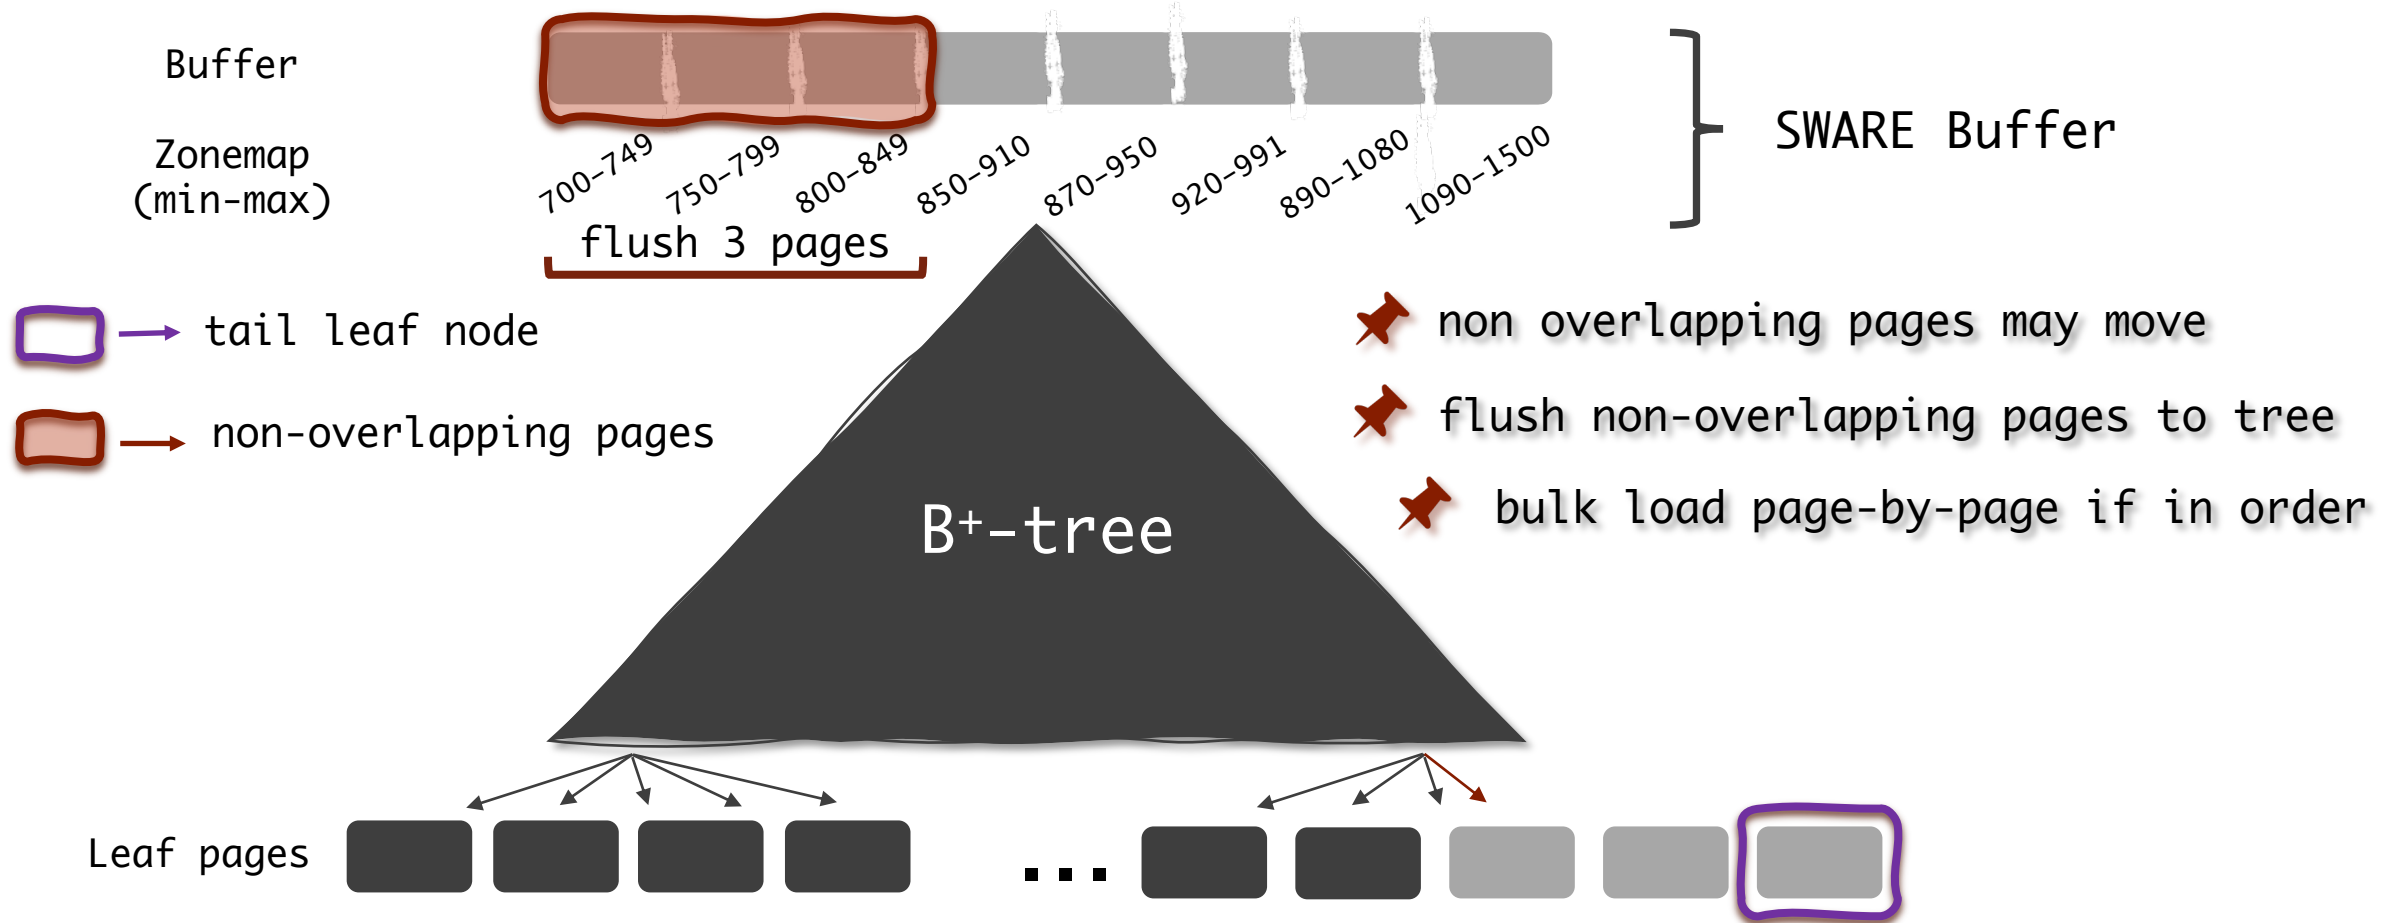


SWARE framework can be applied to any tree-index!


SWARE Ingestions



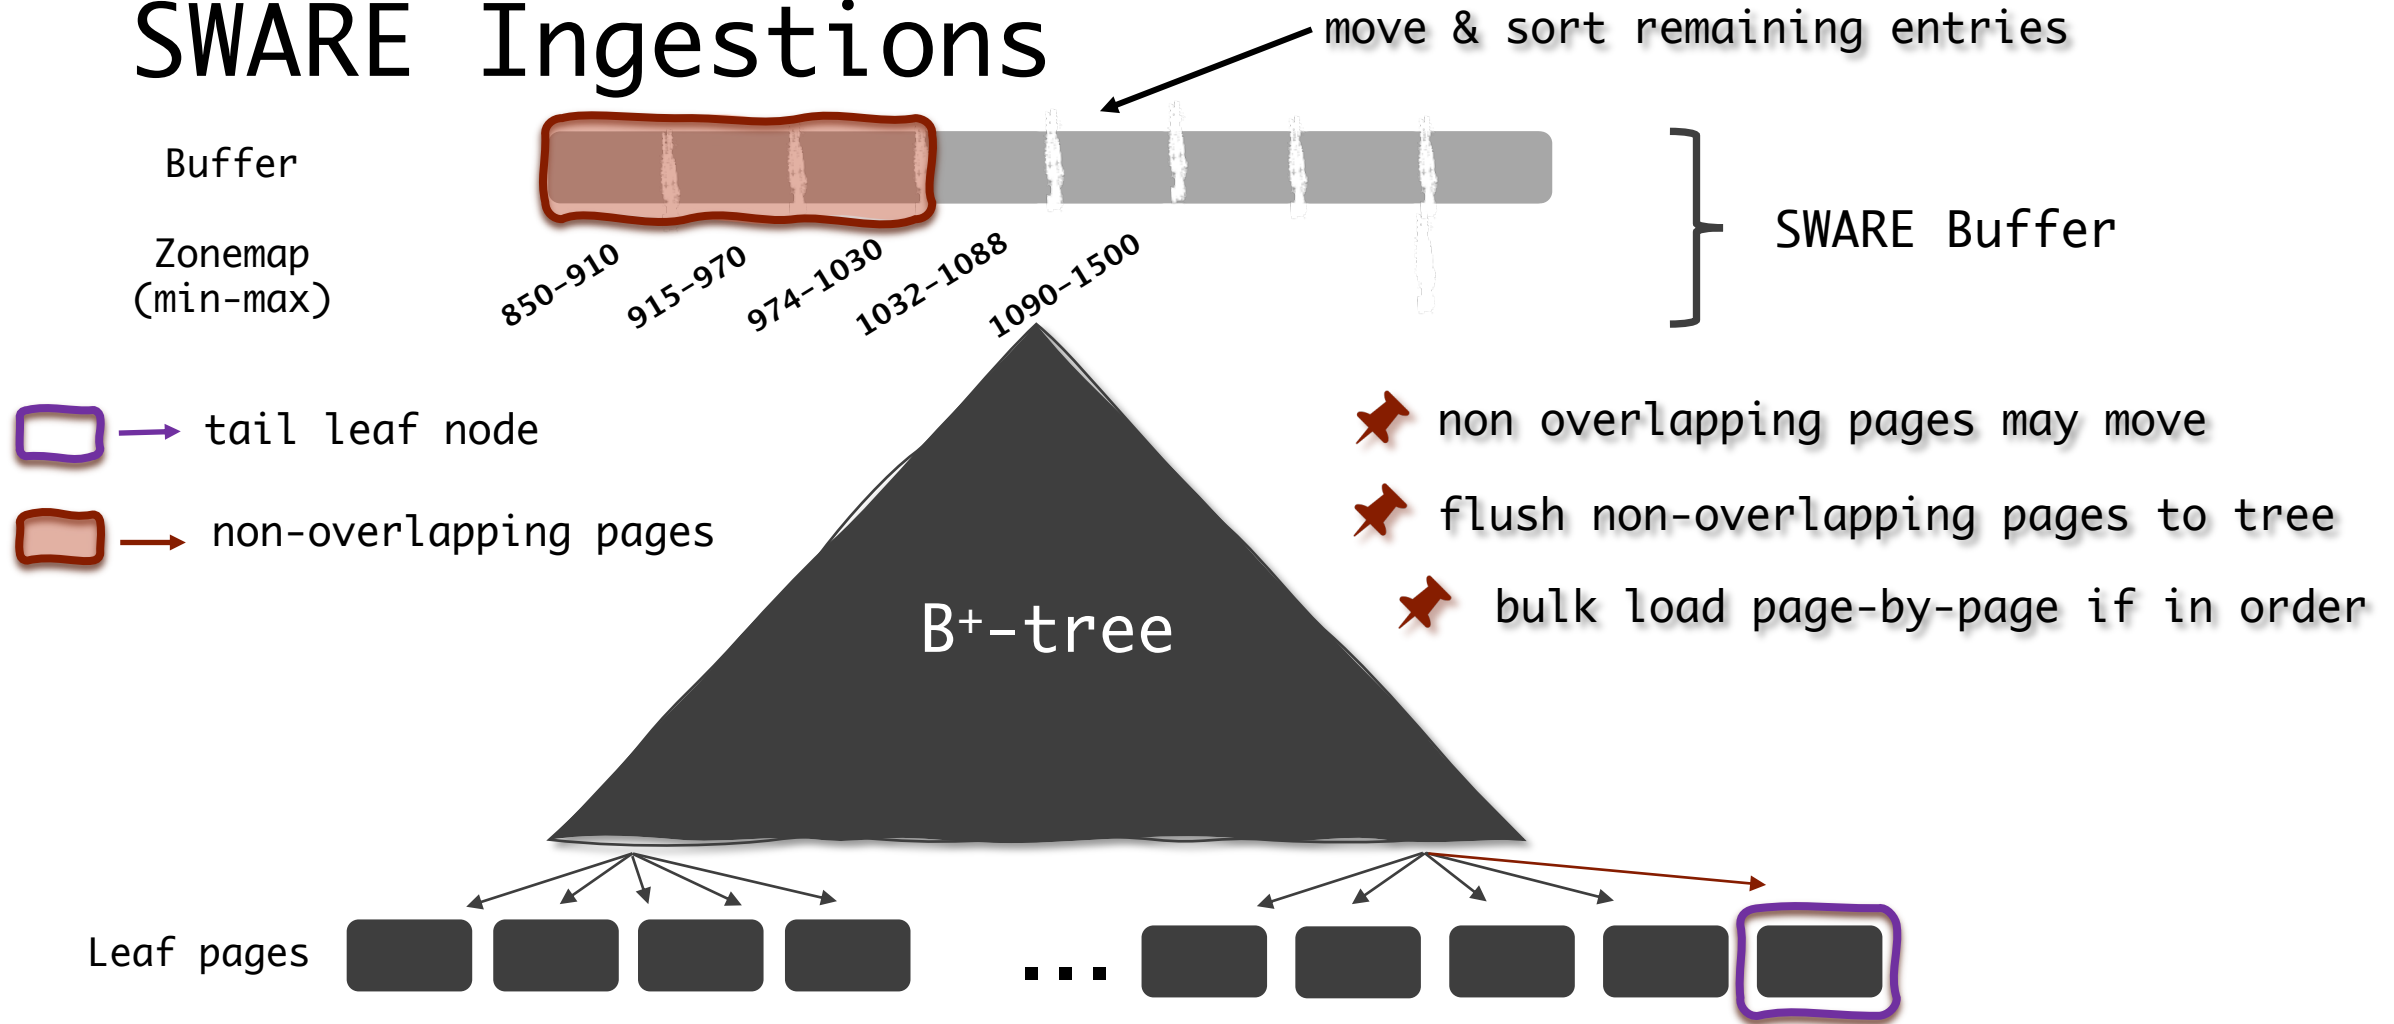
SWARE Ingestions




SWARE Ingestions




SWARE Ingestions



SWARE Ingestions

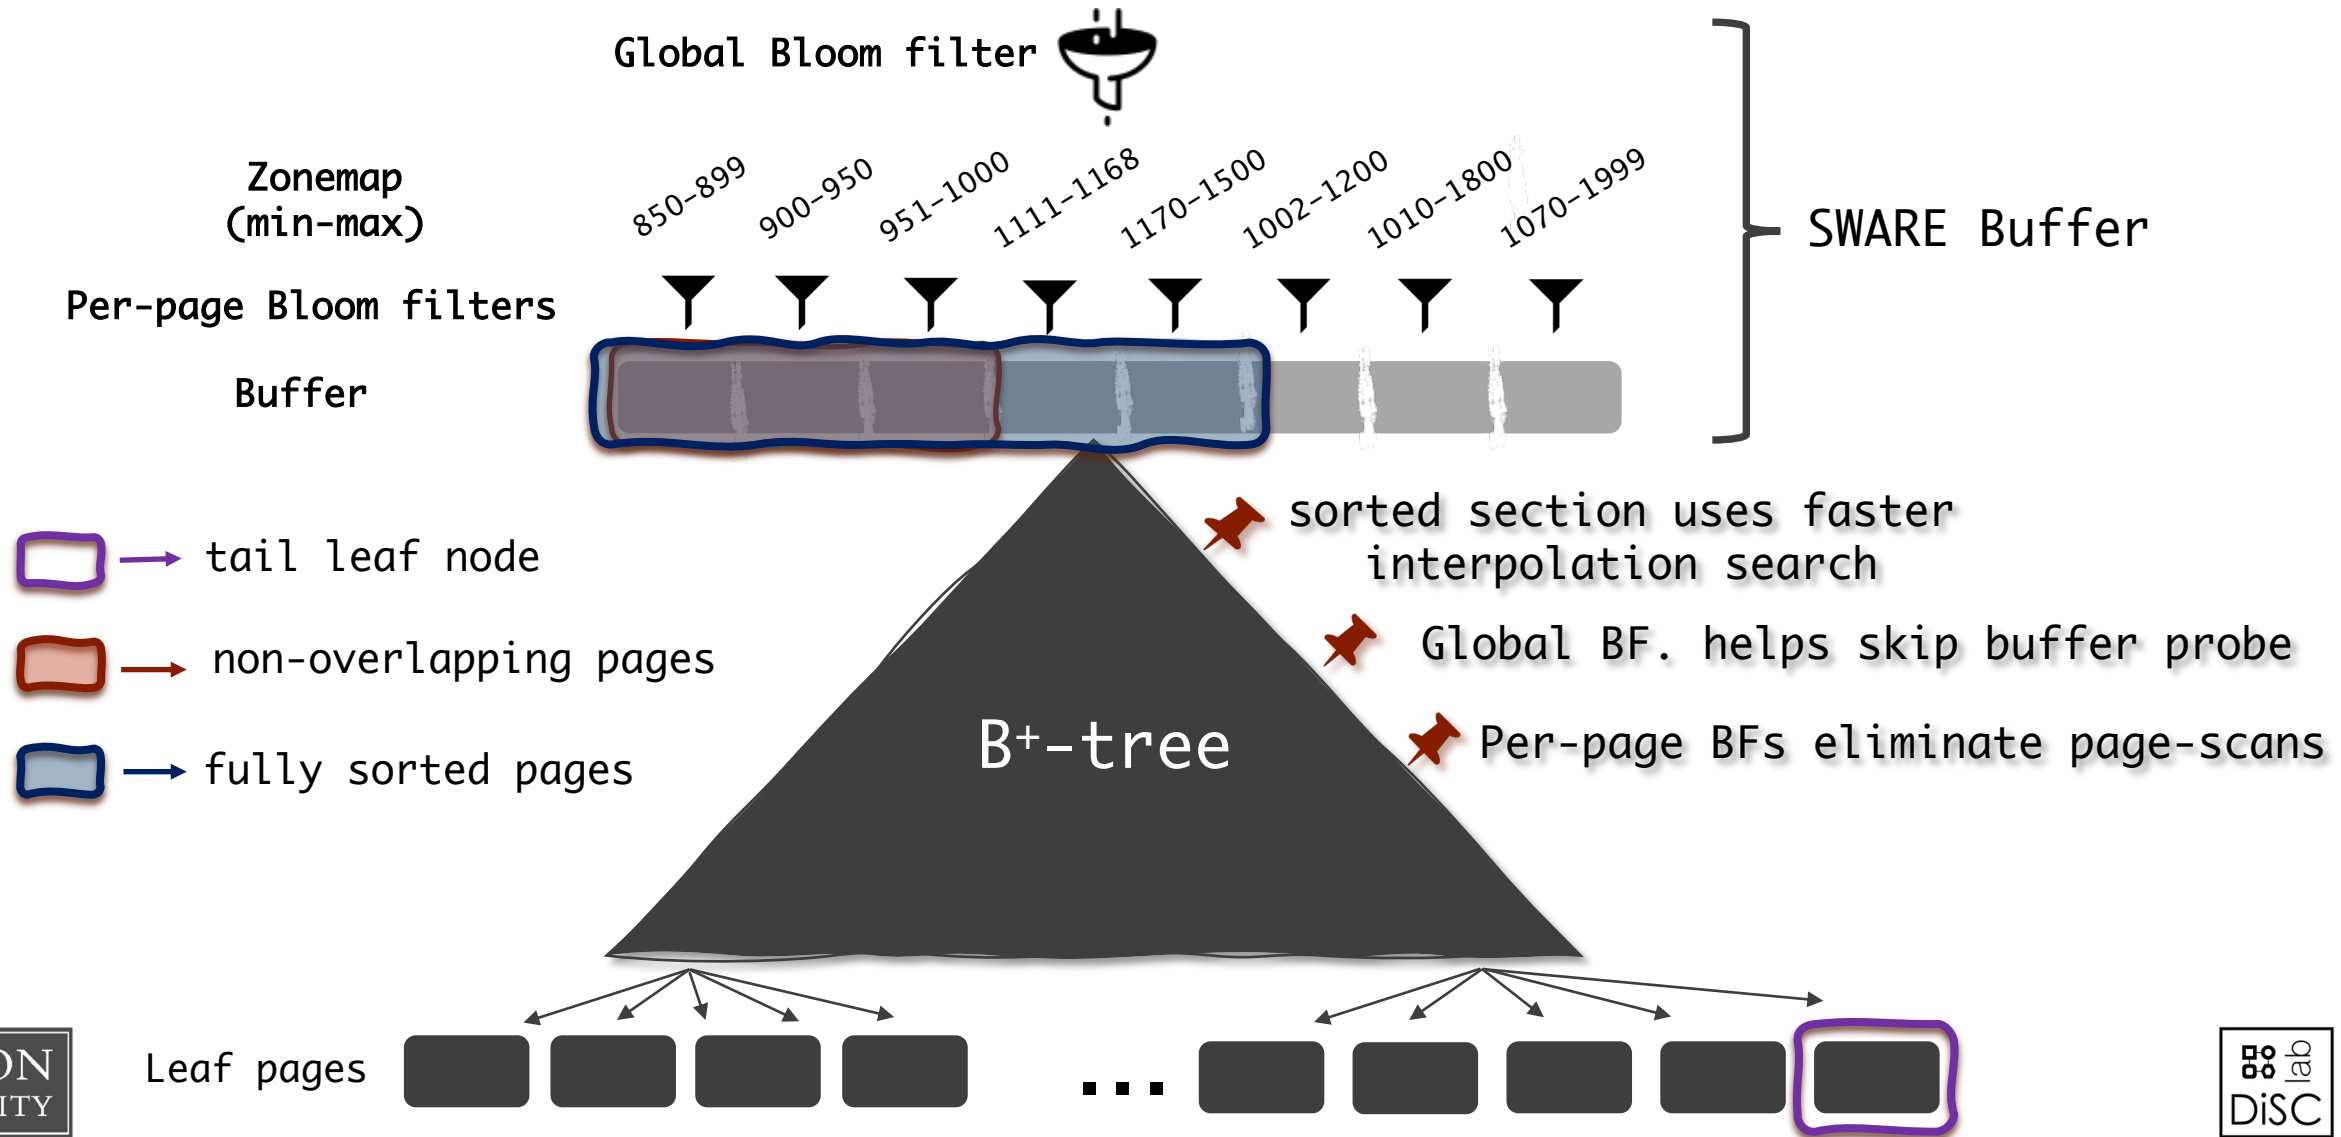


SWARE Ingestions



How do lookups work?

Overall Structure for Queries




How do we evaluate SWARE?

Benchmark on Data Sortedness (TPCTC 2022)

#. unordered entries = K


[BenMoshe, ICDT 2011]



max. displacement among unordered entries = L



Insert Only



Mixed Workloads
(interleaved reads and writes)

Experimental Setup

Metrics:

1. Overall performance (speedup)
2. Raw performance (latency)

Workload Generator: BoDS

1. 500M Integer keys (~ 4GB)
2. Random existing lookups


System Setup:

1. Intel Xeon Gold 5230
2. 2.1GHZ processor w. 20 cores
3. 384GB RAM, 28MB L3 cache


Default Index Setup:

1. Buffer = 40MB; flush \leq 50%
2. BFs = 10 BPK; Murmur Hash
3. Split at 80%


Overall Performance




Overall Performance



Overall Performance




Raw Ingestion Performance




ingestion latency reduced between 27-90%

Raw Ingestion Performance



ingestion latency reduced between 27-90%



bulk loading is maximized with high data sortedness

Space Efficiency

Sortedness Degree	#. Nodes (#. Internal, #. Leaf)	
	B+ tree	SA B+ tree
Fully Sorted	2.004M (8K, 1.996M)	0.52x
Near-Sorted	1.847M (7K, 1.840M)	0.6x
Less-Sorted	1.878M (4.3K, 1.873M)	1.01x

Space Efficiency

Sortedness Degree	#. Nodes (#. Internal, #. Leaf)	
	B+ tree	SA B+ tree
Fully Sorted	2.004M (8K, 1.996M)	0.52x
Near-Sorted	1.847M (7K, 1.840M)	0.6x
Less-Sorted	1.878M (4.3K, 1.873M)	1.01x


increased fill/split factor
helps reduce memory footprint

Summary

Identify “sortedness” as a resource

Smart buffering + bulk index appends
= faster ingestion

Works well with write-heavy or mixed
read-write workloads



Thank You!



: aneeshr@bu.edu



: ramananeesh.com/